

SHARC instruction set



- SHARC programming model.
- SHARC assembly language.
- SHARC memory organization.
- SHARC data operations.
- SHARC flow of control.

SHARC programming model



- Register files:
 - R0-R15 (aliased as F0-F15 for floating point)
- Status registers.
- Loop registers.
- Data address generator registers.
- Interrupt registers.

SHARC assembly language


- Algebraic notation terminated by semicolon:

```
R1=DM (M0, I0) , R2=PM (M8, I8) ; ! comment  
label: R3=R1+R2;
```

data memory access

program memory access

SHARC memory space



SHARC data types



- 32-bit IEEE single-precision floating-point.
- 40-bit IEEE extended-precision floating-point.
- 32-bit integers.
- Memory organized internally as 32-bit words.

SHARC microarchitecture



- Modified Harvard architecture.
 - Program memory can be used to store some data.
- Register file connects to:
 - multiplier
 - shifter;
 - ALU.

SHARC mode registers



- Most important:
 - ASTAT: arithmetic status.
 - STKY: sticky.
 - MODE 1: mode 1.

Rounding and saturation



- Floating-point can be:
 - rounded toward zero;
 - rounded toward nearest.
- ALU supports saturation arithmetic (ALUSAT bit in MODE1).
 - Overflow results in max value, not rollover.

Multiplier



Fixed-point operations can accumulate into local MR registers or be written to register file. Fixed-point result is 80 bits.

Floating-point results always go to register file.

Status bits: negative, under/overflow, invalid, fixed-point underflow, floating-point underflow, floating-point invalid.

ALU/shifter status flags



ALU:

- zero, overflow, negative, fixed-point carry, inputsign, floating-point invalid, last op was floating-point, compare accumulation registers, floating-point under/overflow, fixed-point overflow, floating-point invalid

Shifter:

- zero, overflow, sign

Flag operations



- All ALU operations set AZ (zero), AN (negative), AV (overflow), AC (fixed-point carry), AI (floating-point invalid) bits in ASTAT.
- STKY is sticky version of some ASTAT bits.

Example: data operations

- Fixed-point $-1 + 1 = 0$:
 - $AZ = 1, AU = 0, AN = 0, AV = 0, AC = 1, AI = 0$.
 - STKY bit AOS (fixed point underflow) not set.
- Fixed-point $-2*3$:
 - $MN = 1, MV = 0, MU = 1, MI = 0$.
 - Four STKY bits, none of them set.
- LSHIFT 0x7fffffff BY 3: $SZ=0, SV=1, SS=0$.

Multifunction computations



Can issue some computations in parallel:

- dual add-subtract;
- fixed-point multiply/accumulate and add, subtract, average
- floating-point multiply and ALU operation
- multiplication and dual add/subtract

Multiplier operand from R0-R7, ALU operand from R8-R15.

SHARC load/store



- Load/store architecture: no memory-direct operations.
- Two data address generators (DAGs):
 - program memory;
 - data memory.
- Must set up DAG registers to control loads/stores.

DAG1 registers

I0
I1
I2
I3

M0
M1
M2
M3

L0
L1
L2
L3

B0
B1
B2
B3

I4
I5
I6
I7

M4
M5
M6
M7

L4
L5
L6
L7

B4
B5
B6
B7

Data address generators



Provide indexed, modulo, bit-reverse indexing.

MODE1 bits determine whether primary or alternate registers are active.

BASIC addressing



- Immediate value:

`R0 = DM(0x20000000);`

- Direct load:

`R0 = DM(_a);` ! Loads contents of `_a`

- Direct store:

`DM(_a) = R0;` ! Stores `R0` at `_a`

Post-modify with update

- I register holds base address.
- M register/immediate holds modifier value.

`R0 = DM(I3, M3) ! Load`

`DM(I2, 1) = R1 ! Store`

- Circular buffer: L register is buffer start index, B is buffer base address.

Data in program memory

- Can put data in program memory to read two values per cycle:

$F0 = DM(M0, I0), F1 = PM(M8, I9);$

- Compiler allows programmer to control which memory values are stored in.

Example: C assignments

- C:

```
x = (a + b) - c;
```

- Assembler:

```
R0 = DM(_a) ! Load a
```

```
R1 = DM(_b); ! Load b
```

```
R3 = R0 + R1;
```

```
R2 = DM(_c); ! Load c
```

```
R3 = R3 - R2;
```

```
DM(_x) = R3; ! Store result in x
```

Example, cont'd.



- C:

`y = a * (b + c);`

- Assembler:

`R1 = DM(_b) ! Load b`

`R2 = DM(_c); ! Load c`

`R2 = R1 + R2;`

`R0 = DM(_a); ! Load a`

`R2 = R2 * R0;`

`DM(_y) = R2; ! Store result in y`

Example, cont'd.

- Shorter version using pointers:

! Load b, c

R2=DM(I1,M5), R1=PM(I8,M13);

R0 = R2+R1, R12=DM(I0,M5);

R6 = R12*R0 (SSI);

DM(I0,M5)=R8; ! Store in y

Example, cont'd.

- C:

```
z = (a << 2) | (b & 15);
```

- Assembler:

```
R0=DM(_a); ! Load a
```

```
R0=LSHIFT R0 by #2; ! Left shift
```

```
R1=DM(_b); R3=#15; ! Load immediate
```

```
R1=R1 AND R3;
```

```
R0 = R1 OR R0;
```

```
DM(_z) = R0;
```

SHARC program sequencer



Features:

- instruction cache;
- PC stack;
- status registers;
- loop logic;
- data address generator;

Conditional instructions



Instructions may be executed conditionally.

Conditions come from:

- arithmetic status (ASTAT);
- mode control 1 (MODE1);
- flag inputs;
- loop counter.

SHARC jump



- Unconditional flow of control change:

`JUMP foo`

- Three addressing modes:
 - direct;
 - indirect;
 - PC-relative.

Branches



Types: CALL, JUMP, RTS, RTI.

Can be conditional.

Address can be direct, indirect, PC-relative.

Can be delayed or non-delayed.

JUMP causes automatic loop abort.

Example: C if statement

- C:

```
if (a > b) { x = 5; y = c + d; }  
else x = c - d;
```

- Assembler:

! Test

```
R0 = DM(_a); R1 = DM(_b);
```

```
COMP(R0,R1); ! Compare
```

```
IF GE JUMP fblock;
```

C if statement, cont'd.



! True block

tblock: R0 = 5; ! Get value for x

DM(_x) = R0;

R0 = DM(_c); R1 = DM(_d);

R1 = R0+R1;

DM(_y) = R1;

JUMP other; ! Skip false block

C if statement, cont'd.



! False block

```
fblock: R0 = DM(_c);
```

```
    R1 = DM(_d);
```

```
    R1 = R0-R1;
```

```
    DM(_x) = R1;
```

```
other:  ! Code after if
```

Fancy if implementation



- C:

```
if (a>b) y = c-d; else y = c+d;
```

- Use parallelism to speed it up---compute both cases, then choose which one to store.

Fancy if implementation, cont'd.



! Load values

```
R1=DM(_a); R2=DM(_b);
```

```
R3=DM(_c); R4=DM(_d);
```

! Compute both sum and difference

```
R12 = r2+r4, r0 = r2-r4;
```

! Choose which one to save

```
comp(r8,r1);
```

```
if ge r0=r12;
```

```
dm(_y) = r0 ! Write to y
```


DO UNTIL loops

DO UNTIL instruction provides efficient looping:

```
LCNTR=30, DO label UNTIL LCE;  
R0=DM(I0,M0), F2=PM(I8,M8);  
R1=R0-R15;  
label: F4=F2+F3;
```

Diagram illustrating the components of a DO UNTIL loop:

- Loop length:** Indicated by an arrow pointing from the label `label:` to the start of the loop body.
- Last instruction in loop:** Indicated by an arrow pointing from the text "Last instruction in loop" to the instruction `F4=F2+F3;`.
- Termination condition:** Indicated by an arrow pointing from the text "Termination condition" to the instruction `LCE;` in the `DO label UNTIL LCE;` statement.

Example: FIR filter



- C:

```
for (i=0, f=0; i<N; i++)  
    f = f + c[i]*x[i];
```

FIR filter assembler



```
! setup
  I0=_a; I8=_b; ! a[0] (DAG0), b[0] (DAG1)
  M0=1; M8=1 ! Set up increments
! Loop body
  LCNTR=N, DO loopend UNTIL LCE;
  ! Use postincrement mode
  R1=DM(I0,M0), R2=PM(I8,M8);
  R8=R1*R2;
loopend: R12=R12+R8;
```

Optimized FIR filter code

```
I4=_a; I12=_b;  
R4 = R4 xor R4, R1=DM(I4,M6),  
R2=PM(I12,M14);  
MR0F = R4, MODIFY(I7,M7);  
! Start loop  
LCNTR=20, DO(PC,loop) UNTIL LCE;  
loop: MR0F=MR0F+42*R1 (SSI), R1=DM(I4,M6),  
R2=PM(I12,M14);  
! Loop cleanup  
R0=MR0F;
```

SHARC subroutine calls



- Use CALL instruction:

```
CALL foo;
```

- Can use absolute, indirect, PC-relative addressing modes.
- Return using RTS instruction.

PC stack



PC stack: 30 locations X 24 instructions.

Return addresses for subroutines, interrupt service routines, loops held in PC stack.

Example: C function



- C:

```
void f1(int a) { f2(a); }
```

- Assembler:

```
f1:  R0=DM(I1,-1); ! Load arg into R0
     DM(I1,M1)=R0; ! Push f2's arg
     CALL f2;
     MODIFY(I1,-1); ! Pop element
     RTS;
```

Important programming reminders



- Non-delayed branches (JUMP, CALL, RTS, RTI) do not execute 2 following instructions. Delayed branches are available.
- Cache miss costs at least one cycle to allow program memory bus to complete.

Important programming reminders, cont'd



- Extra cache misses in loops:
 - misses on first and last loop iteration if data memory is accessed in last 2 instructions of loop;
 - 3 misses if loop has only one instruction which requires a program memory bus access.
- 1-instr. loops should be executed 3 times, 2-instr. loops 2 times to avoid NOPs.

Important programming reminders, cont'd



- NOPs added for DAG register write followed by DAG data addressing in same register bank.
- Can program fixed wait states or ACK.
- Interrupt does not occur until 2 instructions after delayed branch.
- Initialize circular buffer by setting L to positive value, loading B to base.

Important programming reminders, cont'd



- Some DAG register transfers are disallowed.
- When given 2 writes to same register file in same cycle, only one actually occurs.
- Fixed- to floating-point conversion always rounds to 40 bits.
- Only DM bus can access all memory spaces.

Important programming reminders, cont'd



- When mixing 32-bit and 48-bit words in a block, all instructions must be below data.
- 16-bit short words are extended to 32 bits.
- For dual data access, use DM for data-only access, PM for mixed data/instruction block. Instruction comes from cache.
- A variety of conditions cause stalls.